

Table of Contents

Lab 02: Gradient Descent

1. Loss landscape
2. The “Gradient” in Gradient Descent

3. Forward & Backward

- 3.1. Forward
- 3.2. Backward
4. Implementation
 - 4.1. Import library
 - 4.2. Create data
 - 4.3. Training

Lab 02: Gradient Descent

Copyright © Department of Computer Science, University of Science, Vietnam National University, Ho Chi Minh City

- Student name: Âu Dương Khang
- ID: 21127621

How to do your homework

- You will work directly on this notebook; the word **TODO** indicates the parts you need to do.
- You can discuss the ideas as well as refer to the documents, but *the code and work must be yours*.

How to submit your homework

- Before submitting, save this file as <ID>.jl . For example, if your ID is 123456, then your file will be 123456.jl . And export to PDF with name 123456.pdf then submit zipped source code and pdf into 123456.zip onto Moodle.

Note

Note that you will get 0 point for the wrong submit.

Content of the assignment:

- Gradient Descent

1. Loss landscape




Figure 1. Loss landscape visualized as a 2D plot. Source: [codecamp.vn](https://www.codecamp.vn)

The gradient descent method is an iterative optimization algorithm that operates over a loss landscape (also called an optimization surface). As we can see, our loss landscape has many peaks and valleys based on which values our parameters take on. Each peak is a local maximum that represents very high regions of loss – the local maximum with the largest loss across the entire loss landscape is the global maximum. Similarly, we also have local minimum which represents many small regions of loss. The local minimum with the smallest loss across the loss landscape is our global minimum. In an ideal world, we would like to find this global minimum, ensuring our parameters take on the most optimal possible values.

Each position along the surface corresponds to a particular loss value given a set of parameters \mathbf{W} (weight matrix) and \mathbf{b} (bias vector). Our goal is to try different values of \mathbf{W} and \mathbf{b} , evaluate their loss, and then take a step towards more optimal values that (ideally) have lower loss.

2. The “Gradient” in Gradient Descent

We can use \mathbf{W} and \mathbf{b} and to compute a loss function L or we are able to find our relative position on the loss landscape, but **which direction** we should take a step to move closer to the minimum.

- All We need to do is follow the slope of the gradient $\nabla_{\mathbf{W}}$. We can compute the gradient $\nabla_{\mathbf{W}}$ across all dimensions using the following equation:

$$\frac{df(x)}{dx} = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

- But, this equation has 2 problems:
 - 1. It’s an **approximation** to the gradient.
 - 2. It’s painfully slow.

In practice, we use the **analytic gradient** instead.

3. Forward & Backward

In this section, you will be asked to fill in the black to form the forward process and backward process with the data defined as follows:

- Feature: \mathbf{X} (shape: $n \times d$, be already used bias trick)
- Label: \mathbf{y} (shape: $n \times 1$)
- Weight: \mathbf{W} (shape: $d \times 1$)

3.1. Forward

TODO: Consider one sample \mathbf{x}_i . Fill in the blank

$$h_i = \mathbf{x}_i^T W \Rightarrow \frac{\partial h_i}{\partial W} = x_i$$

$$\hat{y}_i = \sigma(h_i) \Rightarrow \frac{\partial \hat{y}_i}{\partial h_i} = \sigma(h_i) \times (1 - \sigma(h_i))$$

$$loss_i = (\hat{y}_i - y_i)^2 \Rightarrow \frac{\partial loss_i}{\partial \hat{y}_i} = 2(\hat{y}_i - y_i)$$

- `md"""`
- `**TODO**:` Consider one sample \mathbf{x}_i . Fill in the blank
-
- `$$h_i = \mathbf{x}_i^T W \Rightarrow \frac{\partial h_i}{\partial W} = x_i$$`
-
- `$$\hat{y}_i = \sigma(h_i) \Rightarrow \frac{\partial \hat{y}_i}{\partial h_i} = \sigma(h_i) \times (1 - \sigma(h_i))$$`
-
- `$$loss_i = (\hat{y}_i - y_i)^2 \Rightarrow \frac{\partial loss_i}{\partial \hat{y}_i} = 2(\hat{y}_i - y_i)$$`
-
- `"""`

3.2. Backward

Our loss function is MSE:

$$Loss = \frac{1}{n} \sum_{i=1}^n loss_i = \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

Goal: Compute $\nabla Loss = \frac{\partial Loss(W)}{\partial W}$

How to compute $\nabla Loss$?: Use Chain-rule. Your work is to fill in the blank

TODO: Fill in the blank

$$\begin{aligned} \nabla Loss &= \frac{\partial Loss(W)}{\partial W} \\ &= \frac{1}{n} \sum_{i=1}^n \frac{\partial Loss_i}{\partial \hat{y}_i} \times \frac{\partial \hat{y}_i}{\partial W} \\ &= \frac{1}{n} \sum_{i=1}^n 2(\hat{y}_i - y_i) \times x_i \end{aligned}$$

```
• md"""
• Our loss function is MSE:
•
• $$Loss = \frac{1}{n} \sum_{i=1}^n loss_i = \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$$
•
• **Goal**: Compute $\nabla Loss = \frac{\partial Loss(W)}{\partial W}$
•
• **How to compute $\nabla Loss$?**: Use Chain-rule. Your work is to fill in the blank
•
• **TODO**: Fill in the blank
•
• $$\nabla Loss = \frac{\partial Loss(W)}{\partial W}$
• $$= \frac{1}{n} \sum_{i=1}^n \frac{\partial Loss_i}{\partial \hat{y}_i} \times \frac{\partial \hat{y}_i}{\partial W}$
•
• $$= \frac{1}{n} \sum_{i=1}^n 2(\hat{y}_i - y_i) \times x_i$
• """
```

4. Implementation


4.1. Import library

```
• using Distributions , Plots , LinearAlgebra , Random
```

Cell deleted (UNDO)

- [Random](#).seed!(2024)

4.2. Create data



```

• begin
•     # DOT NOT MODIFY THIS CODE
•     # generate a 2-class classification problem with 1,000 data points, each data
•     # point is a 2D feature vector
•     # number of data points
•     n = 1000
•
•     # dimensionality of data
•     d = 2
•
•     # mean
•      $\mu$  = 5
•
•     # variance
•      $\Sigma$  = 8
•
•     # Generate two class for synthesized data
•     positive = rand(MvNormal( $[\Sigma, \mu]$ , 3 .* [1 ( $\mu - d$ )/ $\mu$ ; ( $\mu - d$ )/ $\mu$  d])), n ÷ 2)
•     negative = rand(MvNormal( $[-\mu, \Sigma]$ , 3 .* [1 ( $\mu - d$ )/ $\mu$ ; ( $\mu - d$ )/ $\mu$  d])), n ÷ 2)
•
•     # Combine two class of generated data.
•     # X = features
•     # y = label
•     X = hcat(positive, negative)
•     y = vcat(ones(n ÷ 2) .- 1, ones(n ÷ 2))'
•
•     # Visualization
•     plt = scatter(positive[1, :], positive[2, :], label="y = 1")
•     scatter!(plt, negative[1, :], negative[2, :], label="y = -1")
•     # DOT NOT MODIFY THIS CODE

```

Cell deleted (UNDO)

4x1000 Matrix{Float64}:

```
7.00921  5.40614  7.77031  10.2313  ...  -0.235459  -2.19833  -4.99492  -5.55589
4.07511  1.63066  9.29032  8.68727  ...  11.5486   10.9082   7.40612   8.02935
1.0      1.0      1.0      1.0      ...  1.0      1.0      1.0      1.0
0.0      0.0      0.0      0.0      ...  1.0      1.0      1.0      1.0
```

```
• begin
•   # insert a column of 1's as the last entry in the feature matrix
•   # -- allows us to treat the bias as a trainable parameter
•
•   X_aug = vcat(X, ones(n)')
•   data = vcat(X_aug, y)
• end
```

```
(700x3 Matrix{Float64}: , [0.0, 1.0, 1.0, 1.0, 1.0, 0.0, 0.0, 1.0, 1.0,      more ,0.0], 300
 7.05549  4.70516  1.0
-6.87358  3.98823  1.0
-2.05821  6.87055  1.0
-5.71526  7.10931  1.0
-7.14349  3.62296  1.0
 8.85531  6.33639  1.0
10.045    5.23048  1.0
⋮
-5.93771  6.71143  1.0
-5.38006  8.4446   1.0
-4.582    10.4695  1.0
-4.57232  7.8681   1.0
 9.01379  5.11398  1.0
 7.12285  3.71975  1.0
```

```
• begin
•   # DOT NOT MODIFY THIS CODE
•   # Split data, use 50% of the data for training and the remaining 50% for testing
•   # Prepare data
•   D = data'[shuffle(1:end), :]
•
•   # Calculate the number of samples for each split
•   n_train = Int(n * 0.7)
•
•   # Split the samples into train, and test sets
•   train_data = D[begin:n_train, :]
•   test_data = D[n_train + 1: end, :]
•   println(size(train_data), size(test_data))
•
•   # Move samples to train-test features and labels
•   X_train, y_train, X_test, y_test = train_data[:,1:3], train_data[:,4],
•   test_data[:,1:3], test_data[:,4]
•   # DOT NOT MODIFY THIS CODE
• end
```

(700, 4)(300, 4)



4.3. Training

Sigmoid function and derivative of the sigmoid function

sigmoid_deriv (generic function with 1 method)

```
• begin
•   function sigmoid_activation(x)
•       #TODO
•       """compute the sigmoid activation value for a given input"""
•       #return?
•       return 1.0 ./ (1.0 .+ exp.(-x))
•
•   end
•
•   function sigmoid_deriv(x)
•       #TODO
•       """
•       Compute the derivative of the sigmoid function ASSUMING
•       that the input 'x' has already been passed through the sigmoid
•       activation function
•       """
•       #return?
•       return sigmoid_activation(x) .* (1.0 .- sigmoid_activation(x))
•
•   end
• end
```

Compute output

predict (generic function with 1 method)

```
• begin
•   function compute_h(W, X)
•       #TODO
•       """
•       Compute output: Take the inner product between our features 'X' and the weight
•       matrix 'W'
•       """
•       # return?
•       return X*W
•
•   end
•
•   function predict(W, X)
•       #TODO
•       """
•       Take the inner product between our features and weight matrix,
•       then pass this value through our sigmoid activation
•       """
•       # preds = ...
•       preds = sigmoid_activation(compute_h(W,X))
•
•       # apply a step function to threshold the outputs to binary
•       # class labels
•       preds[preds .<= 0.5] .= 0
•       preds[preds .> 0] .= 1
•
•       return preds
•   end
• end
```

Compute gradient

compute_gradient (generic function with 1 method)

```
• begin
•   function compute_gradient(error, y_hat, trainX)
•       #TODO
•       """
•       the gradient descent update is the dot product between our
•       features and the error of the sigmoid derivative of
•       our predictions
•       """
•       dactivation = sigmoid_deriv(y_hat)
•       gradient = trainX' * (error .* dactivation)
•       # return?
•       return gradient
•
•   end
• end
```

Cell deleted ([UNDO](#))

ction

train (generic function with 1 method)

```
• begin
•     function train(W, trainX, trainY, learning_rate, num_epochs)
•         losses = []
•         for epoch in 1:num_epochs
•             y_hat = sigmoid_activation(compute_h(W, trainX))
•             # now that we have our predictions, we need to determine the
•             # 'error', which is the difference between our predictions and
•             # the true values
•             error = y_hat - trainY
•             append!(losses, 0.5 * sum(error .^ 2))
•             grad = compute_gradient(error, y_hat, trainX)
•             W -= learning_rate * grad
•
•             if epoch == 1 || epoch % 5 == 0
•                 println("Epoch=$epoch; Loss=$(losses[end])")
•             end
•         end
•         return W, losses
•     end
• end
```


Initialize our weight matrix and list of losses

```
• md"""
• ##### Initialize our weight matrix and list of losses
• """
```

0.1

```
• begin
•     #initialize our weight matrix and necessary hyperparameters
•     W = rand(Normal(), (size(X_train)[2], 1))
•     num_epochs=100
•     learning_rate=0.1
• end
```

Train our model



```

• begin
•     #training model
•      $\theta$ , losses = train(W, X_train, y_train, learning_rate, num_epochs)
•     #visualiza training process
•     plot(1:num_epochs, losses, legend=false)
• end

```

```

Epoch=1; Loss=162.89330008883746
Epoch=5; Loss=0.4996629372262392
Epoch=10; Loss=0.213901725493055
Epoch=15; Loss=0.00519495524271303
Epoch=20; Loss=0.0013469312005308456
Epoch=25; Loss=0.000604331216694869
Epoch=30; Loss=0.0003413439086842821
Epoch=35; Loss=0.00021895307403778155
Epoch=40; Loss=0.00015226302432830107
Epoch=45; Loss=0.00011197159288541358
Epoch=50; Loss=8.578542214448916e-5
Epoch=55; Loss=6.78146276933873e-5
Epoch=60; Loss=5.495070756475468e-5
Epoch=65; Loss=4.5427387866039997e-5
Epoch=70; Loss=3.818080640171032e-5
Epoch=75; Loss=3.253916550707599e-5
Epoch=80; Loss=2.8061322021449194e-5
Epoch=85; Loss=2.4447897715880722e-5
Epoch=90; Loss=2.1489930265636757e-5
Epoch=95; Loss=1.9037981612168296e-5
Epoch=100; Loss=1.6982893206780018e-5

```

Evaluate result

Cell deleted (UNDO)

```

• begin
•   y_pred = predict(W, X_test)
•   true_positives = 0
•   false_positives = 0
•   true_negatives = 0
•   false_negatives = 0
•
•   # Calculate true positives, false positives, false negatives, and true negatives
•   for (true_label, predicted_label) in zip(y_test, y_pred)
•     if true_label == 1 && predicted_label == 1
•       true_positives += 1
•     elseif true_label == 0 && predicted_label == 1
•       false_positives += 1
•     elseif true_label == 1 && predicted_label == 0
•       false_negatives += 1
•     elseif true_label == 0 && predicted_label == 0
•       true_negatives += 1
•     end
•   end
•
•   # Calculate precision, recall, and F1-score
•   accuracy = (true_positives + true_negatives) / (true_positives + false_positives
+ true_negatives + false_negatives)
•   precision = true_positives / (true_positives + false_positives)
•   recall = true_positives / (true_positives + false_negatives)
•   f1_score = 2 * precision * recall / (precision + recall)
•
•   # Display
•   print("acc: $accuracy, precision: $precision, recall: $recall, f1_score:
$f1_score\n")
• end


```

```

acc: 0.52, precision: 1.0, recall: 0.03355704697986577, f1_score: 0.0649
3506493506493

```





```

• begin
•   # Create a scatter plot
•   scatter(X_test'[1,:][y_test == 0], X_test'[2,:][y_test == 0], label="Class 0",
•           color=:blue, legend=:topright, markersize=4)
•   scatter!(X_test'[1,:][y_test == 1], X_test'[2,:][y_test == 1], label="Class 1",
•            color=:red, markersize=4)
•
•   # Getting decision boundary configuration
•   b = θ[3]
•   θm1 = θ[1:2]
•
•   decision(x) = θm1' * x + b
•
•   D_test = ([
•       tuple.(eachcol(hcat(X_test'[1,:][y_test == 0], X_test'[2,:][y_test == 0])),
•       1)
•       tuple.(eachcol(hcat(X_test'[1,:][y_test == 1], X_test'[2,:][y_test == 1])),
•       -1)
•   ])
•
•   # Max, min for visualization decision boundary
•   x_min = minimum(map((p) -> p[1][1], D_test))
•   y_min = minimum(map((p) -> p[1][2], D_test))
•   x_max = maximum(map((p) -> p[1][1], D_test))
•   y_max = maximum(map((p) -> p[1][2], D_test))
•
•   # Display decision boundary
•   contour!(plt, x_min:0.1:x_max, y_min:0.1:y_max,
•            (x, y) -> decision([x, y]),
•            levels=[0], linestyle=:solid, label="Decision boundary",
•            lorbar_entry=false, color=:green)
• end

```

Cell deleted (UNDO)

Cell deleted (UNDO)

TODO: Study about accuracy, recall, precision, f1-score.

- Accuracy: This is the ratio of true positives and true negatives to all cases. It's a measure of how good our model is at predicting the correct category (classes or labels).

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + NP + NN}$$

- Recall: This is the ratio of true positives to all actual positives. Recall is important when the cost of missing a positive instance is high.

$$\text{Recall} = \frac{TP}{TP + FN}$$

- Precision: This is the ratio of true positives to all predicted positives. We focus on precision when we need our predictions to be correct, i.e., ideally, we want to make sure our model is right when it predicts a label.

$$\text{Precision} = \frac{TP}{TP + FP}$$

- F1: This is the weighted average of precision and recall and is usually more useful than accuracy, especially if the class distribution is uneven or the cost of false positives and false negatives are different. F1 score is the harmonic mean of precision and recall and becomes 1 only when both are 1. F1 score can be calculated by $F1 = 2 * (\text{precision} * \text{recall}) / (\text{precision} + \text{recall})$

$$F1 = \frac{2 \times \text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

Looking at the results, it seems like the model has a high precision (1.0), meaning that when it predicts a positive result, it's usually correct. However, its recall is quite low (0.033576497989677), indicating that it's missing a lot of actual positive instances. The F1-score (0.06493506493506493) reflects this imbalance between precision and recall. The accuracy (0.52) isn't very high either, suggesting that the model isn't performing very well overall.

- `md"""`
- `**TODO: Study about accuracy, recall, precision, f1-score.**`
- `- Accuracy: This is the ratio of true positives and true negatives to all cases. It's a measure of how good our model is at predicting the correct category (classes or labels).`
- `$$Accuracy = \frac{TP + TN}{TP + TN + NP + NN}$$`
- `- Recall: This is the ratio of true positives to all actual positives. Recall is important when the cost of missing a positive instance is high.`
- `$$Recall = \frac{TP}{TP + FN}$$`
- `- Precision: This is the ratio of true positives to all predicted positives. We focus on precision when we need our predictions to be correct, i.e., ideally, we want to make sure our model is right when it predicts a label.`

Cell deleted (UNDO) `\frac{TP}{TP + FP}$$`

- F1: This is the weighted average of precision and recall and is usually more useful than accuracy, especially if the class distribution is uneven or the cost of false positives and false negatives are different. F1 score is the harmonic mean of precision and recall and becomes 1 only when both are 1. F1 score can be calculated by $F1 = 2 * (\text{precision} * \text{recall}) / (\text{precision} + \text{recall})$

$$F1 = \frac{2 * \text{Precision} * \text{Recall}}{\text{Precision} + \text{Recall}}$$

Looking at the results, it seems like the model has a high precision (1.0), meaning that when it predicts a positive result, it's usually correct. However, its recall is quite low (0.033576497989677), indicating that it's missing a lot of actual positive instances. The F1-score (0.06493506493506493) reflects this imbalance between precision and recall. The accuracy (0.52) isn't very high either, suggesting that the model isn't performing very well overall.

TODO: Try out different learning rates. Give me your observations

- Learning rate is used to scale the magnitude of parameter updates during gradient descent. The choice of the value for learning rate can impact two things: 1) how fast the algorithm learns and 2) whether the cost function is minimized or not.
- High learning rate: If the learning rate is too high, the algorithm might overshoot the minimum and diverge, resulting in large oscillations. The cost function may end up bouncing back and forth and take a longer time to converge, or in some cases, it might not converge at all. The value of high learning rate is about higher 0.1.
- Optimal learning rate: An optimal learning rate would lead to a steady and quick convergence to the global minimum. The cost function decreases and eventually stabilizes, leading to the lowest possible loss. The value of optimal learning rate is about $0.01 < lr < 0.1$.
- Low learning rate: If the learning rate is too low, the algorithm will eventually reach the minimum but the convergence will be very slow. This means the algorithm will take a very long time to run and will be computationally expensive. The value of optimal learning rate is about lower 0.01.

```
• md"""
• **TODO: Try out different learning rates. Give me your observations**
• - Learning rate is used to scale the magnitude of parameter updates during gradient
  descent. The choice of the value for learning rate can impact two things: 1) how fast
  the algorithm learns and 2) whether the cost function is minimized or not.
• - High learning rate: If the learning rate is too high, the algorithm might overshoot
  the minimum and diverge, resulting in large oscillations. The cost function may end
  up bouncing back and forth and take a longer time to converge, or in some cases, it
  might not converge at all. The value of high learning rate is about higher 0.1.
• - Optimal learning rate: An optimal learning rate would lead to a steady and quick
  convergence to the global minimum. The cost function decreases and eventually
  stabilizes, leading to the lowest possible loss. The value of optimal learning rate
  is about  $0.01 < lr < 0.1$ .
• - Low learning rate: If the learning rate is too low, the algorithm will eventually
  reach the minimum but the convergence will be very slow. This means the algorithm
  will take a very long time to run and will be computationally expensive. The value of
  optimal learning rate is about lower 0.01.
• """
```

Cell deleted (UNDO)